Google 🔍 Ayush Singh - distributed systems engineer ✕ ⚙ ⋮⋮⋮ You

All | Experience | Projects | ✦ AI Mode | Skills | More | Tools

ayushsingh.is-a.dev › work-experience › distributed-systems-engineer

✦ AI Overview · Executive Summary

Software Engineer III (L4) at **Google** (Munich) with 4+ years building planet-scale **distributed systems**, **low-latency** C++/Java RPC frameworks, fault-tolerant distributed file systems, and real-time streaming pipelines across Google, Komprise, Accenture, and Tookitaki.

- **Planet-Scale Core Infra:** Defaulted architecture across ~5,200 microservice nodes (Search, YouTube, Maps) & built distributed storage serving 100,000+ services.
- **High-Impact Systems ROI**: Delivered $2.5M+/quarter savings (SLA platform) and 94% streaming latency reduction (850ms → <50ms).
- **Fast Track Record**: Promoted L3 → L4 in 15 months; VP Spot Bonus; 3× Readability (C++, TS, API) & 1040+ PRs reviewed.

Top stats

~5,200	100K+	$2.5M+	94%
Prod microservice nodes running my default design	Google services backed by my distributed FileSystem	Quarterly savings via SLA/SLO platform	Streaming latency cut (850ms → <50ms)

Work Experience · 4+ Years in Distributed Systems, Storage & Backend

Google — Software Engineer III (L4), Distributed Systems & Core Infrastructure

Munich, Germany · Nov 2024 – Present · C++ · Java · Kotlin · gRPC · Distributed Microservices · Edge Routing

- **Core C++ & Java Request Frameworks:** Shipped global request-path processors into **Google's core Java backend framework (powering all Google products)** and **core C++ Search serving engine**, serving **billions of users**.
- **Default Architecture Across ~5,200 Nodes:** Authored the org-wide architecture defaulted across **~5,200 production microservice nodes** (210+ modules), cutting signal-processing complexity by **40%** and retiring **1,300+ legacy LOC**.
- **Cross-Product Telemetry & Batch Compute:** Built parallel stream shadow-diffing across **YouTube, Play, Search & Payments** (100% live cut-over parity) and migrated **7 critical identity pipelines** to distributed batch compute.
- **High-Velocity Regulatory Rollouts:** Led end-to-end **UK Online Safety Act & US AppStore** distributed backend launches across Web, Android & iOS in 4 weeks across 5+ teams (**VP Spot Bonus**).

Google — Software Engineer II → III (L4), Distributed Storage & Observability

Hyderabad, India · Jul 2022 – Nov 2024 · C++ · Go · FUSE · Distributed File Systems · Paxos Consensus · OpenTelemetry

- **Distributed File System (100K+ Services):** Engineered fault-tolerant storage and replication for **Google's internal HDFS-like distributed file system**, optimizing mission-critical binary delivery across **100,000+ internal services**.
- **Overload Protection & Consensus Locks:** Shipped server-side admission control & load-shedding, semaphore RPC concurrency limiting in the **FUSE client**, and migrated distributed snapshot locks off **Paxos lock services**.
- **$2.5M+ Quarterly Savings via SLA/SLO Platform:** Architected a highly available observability & automated **SLA/SLO compliance platform** correlating distributed traces and metrics, saving **$2.5M+ per quarter**.
- **OpenTelemetry Collector & Go Gateway:** Co-architected a trace collector over **gRPC-Web** & encrypted **Go asset gateway**; promoted to **L4 in 15 months**.



Ayush Singh

Distributed Systems Engineer · Google L4

Overview | Experience | Skills

Current: SWE III (L4) at Google
Experience: 4+ Years
Location: Munich, Germany
Email: apparently.ayush@gmail.com
Phone: +49 1735 34 8915
Education: NIT Agartala · B.Tech EE (CGPA: 9.32 / 10)
Code Authority: 3× Google Readability (C++, TS, API) · 740+ Reviews
Honors: VP Spot Bonus · 7× Peer Bonuses · GFG Rank #72

Profiles & Links

LinkedIn | Portfolio | GitHub | Email

Technical skills

Languages
C++ Java Go TypeScript Python Scala Kotlin SQL

Distributed Systems & Storage
gRPC Protobuf Distributed FS (HDFS / FUSE) Spanner Paxos / Raft Consensus Apache Spark Kafka ActiveMQ Snowflake Redis

Cloud, Observability & AI Infra
Kubernetes Docker OpenTelemetry GCP AWS Terraform CI/CD WebGPU / LLM Agents

Komprise — Software Engineer Intern (Distributed Storage Infrastructure)

Bengaluru, India · Jan 2022 – Jun 2022 · Python · Multi-Threading · Distributed Systems Testing · CI/CD

- **Parallel Distributed Test Execution:** Built a multi-threaded automation framework executing in parallel across distributed storage nodes, cutting build-verification runtime by **87.5%** and increasing bug detection coverage by **60%**.

Accenture — Data Engineer Intern (Cloud Data Infrastructure)

India · May 2021 – Jul 2021 · AWS S3 · Snowflake · Distributed ETL · Python · SQL

- **10TB+ High-Volume Cloud ETL:** Designed optimized distributed ETL pipelines transferring **10TB+ datasets from AWS S3 to Snowflake** with real-time analytics automation cutting reporting time by **70%**

Tookitaki — Software Engineer Intern (Real-Time Distributed Streaming)

Singapore · Oct 2020 – Jan 2021 · Apache Spark Structured-Streaming · Scala · ActiveMQ · JMS

- **94% Streaming Latency Reduction (850ms → <50ms):** Built a real-time **Apache Spark Structured-Streaming + ActiveMQ** pipeline with custom connection-pooled JMS clients and **exactly-once fault-tolerant recovery**.

✦ Selected Open-Source & Agentic Systems

1. YTM AI Host — Real-Time (<100ms) AI Radio Jockey & Audio Streaming Engine · GitHub

TypeScript · Python · Gemini API · WebLLM / WebGPU · On-Device Inference · Real-Time Audio Pre-Fetching

- **Sub-100ms Streaming AI Architecture:** Built a modular local streaming server and Chrome extension for **YouTube Music & Spotify**, combining **on-device WebGPU/WebLLM inference**, multi-engine TTS (Kokoro JS, XTTS v2), and intelligent pre-fetching/caching for **<100ms real-time audio transitions**.

2. LinkinJMS (High-Performance Spark Queue Connector) & Multi-Agent Incident Triage

- **LinkinJMS (Scala / Apache Spark / ActiveMQ):** Open-source bidirectional streaming connector achieving **<50ms message latency** with **exactly-once delivery**, AMQP/TLS encryption, and automatic failover.
- **Autonomous Multi-Agent Oncall Triage (Google):** Production Python multi-agent system that fans out parallel diagnostic queries across code changes, rollout logs, and time-series metrics to email root-cause reports before oncallers open the alert.

People also ask

What scale of distributed systems does Ayush build at Google?
His C++/Java middleware runs on the critical request path of Search, YouTube & Maps (~5,200 microservice nodes), and his distributed storage systems serve 100K+ services. ˅

How fast was Ayush promoted at Google?
Promoted L3 → L4 (SWE III) in 15 months with a VP of Engineering Spot Bonus, 7× peer bonuses, and 3× Google Readability certifications (C++, TypeScript, API; 740+ code reviews). ˅

What measurable systems impact has Ayush delivered?
$2.5M+/quarter saved via automated SLA/SLO enforcement, 94% streaming latency reduction (850ms → <50ms), 40% lower RPC complexity, and 87.5% faster distributed test execution. ˅

Where did Ayush study and what are his algorithmic credentials?
NIT Agartala (B.Tech EE, 9.32/10 CGPA). Ranked #1 in Institute & #72 Worldwide on GeeksForGeeks (top 0.07% of 100K+ developers), IASc Fellow, and Founding Google DSC Lead. ˅

People also search for

Distributed File Systems (HDFS/FUSE) | High-Throughput gRPC | Apache Spark Streaming | OpenTelemetry | Paxos & Consensus

Gooooooogle ›

● **Munich, Germany** · apparently.ayush@gmail.com · +49 1735 34 8915 · linkedin.com/in/iu-sh